

CBF Programming Assignment

Overview

In this assignment, you will implement a content-based recommender as a LensKit recommender algorithm. LensKit provides tools to produce recommendations from a user; your task is to implement the logic of the recommender itself.

There are 2 parts to this assignment, implementing two variants of a TF-IDF recommender.

Downloads and Resources

- Project template (on course website)
- LensKit for Teaching website (links to relevant documentation and the LensKit tutorial video)
- JavaDoc for included code

Additionally, you will need:

- Java — download the Java 8 JDK. On Linux, install the OpenJDK ‘devel’ package (you will need the devel package to have the compiler).
- An IDE; I recommend IntelliJ IDEA Community Edition.

Notation

Here’s the mathematical notation we are using:

\vec{u} The user’s vector (in this assignment, the user profile vector).

\vec{i} The item vector.

$I(u)$ The set of items rated by user u .

u_t, i_t User u ’s or item i ’s score for tag t

r_{ui} User u ’s rating for item i .

μ_u The average of user u ’s ratings.

Part 1: TF-IDF Recommender with Unweighted Profiles (85 points)

Start by downloading the project template. This is a Gradle project; you can import it into your IDE directly (IntelliJ users can open the build.gradle file as a project). The code should compile as-is; you can test this by running the `build` Gradle target from your IDE, or running `./gradlew build` at the command line.

There are 3 things you need to implement to complete the first part of the assignment:

Compute item-tag vectors (the model) For this task, you need to modify the model builder (`TFIDFModelBuilder`, your modifications go in the `get()` method) to compute the unit-normalized TF-IDF vector for each movie in the data set. We provide the skeleton of this; TODO comments indicate where you need to implement missing pieces. When this piece is done, the model should contain a mapping of item IDs to TF-IDF vectors, normalized to unit vectors, for each item.

Build user profile for each query user The `UserProfileBuilder` interface defines classes that take a user's history – a list of ratings — and produce a vector representing that user's profile. For part 1, the profile should be the sum of the item-tag vectors of all items the user has rated positively (≥ 3.5 stars); this implementation goes in `ThresholdUserProfileBuilder`.

Generate item scores for each user The heart of the recommendation process in many LensKit recommenders is the score method of the item scorer, in this case `TFIDFItemScorer`. Modify this method to score each item by using cosine similarity: the score for an item is the cosine between that item's tag vector and the user's profile vector. Cosine similarity is defined as follows:

$$\cos(u, i) = \frac{\vec{u} \cdot \vec{i}}{\|\vec{u}\|_2 \|\vec{i}\|_2} = \frac{\sum_t u_t i_t}{\sqrt{\sum_t u_t^2} \sqrt{\sum_t i_t^2}}$$

You can run your program from the command line using Gradle:

```
...  
./gradlew recommendBasic -PuserId=42  
...
```

Try different user IDs.

Example Output for Unweighted User Profile

The following example gives actual outputs for users 42 and 91 in the data set. It was executed using `./gradlew recommendBasic -PuserId=42,91` in a Unix-like console.

```
recommendations for user 42:  
2081 (Little Mermaid, The (1989)): 0.500  
2671 (Notting Hill (1999)): 0.292  
2724 (Runaway Bride (1999)): 0.286  
1265 (Groundhog Day (1993)): 0.241  
58299 (Horton Hears a Who! (2008)): 0.224  
96861 (Taken 2 (2012)): 0.215  
1210 (Star Wars: Episode VI - Return of the Jedi (1983)): 0.187  
1202 (Withnail & I (1987)): 0.127  
26048 (Human Condition II, The (Ningen no joken II) (1959)): 0.110  
260 (Star Wars: Episode IV - A New Hope (1977)): 0.109
```

recommendations for user 91:

4816 (Zoolander (2001)): 0.343
8376 (Napoleon Dynamite (2004)): 0.252
1247 (Graduate, The (1967)): 0.232
4361 (Tootsie (1982)): 0.210
6350 (Laputa: Castle in the Sky (Tenkû no shiro Rapyuta) (1986)): 0.169
2167 (Blade (1998)): 0.161
2858 (American Beauty (1999)): 0.159
6947 (Master and Commander: The Far Side of the World (2003)): 0.158
5459 (Men in Black II (a.k.a. MIIB) (a.k.a. MIB 2) (2002)): 0.158
2724 (Runaway Bride (1999)): 0.151

Part 2: Weighted User Profile (15 points)

For this part, adapt your solution from Part 1 to compute weighted user profiles. Put your weighted user profile code in `WeightedUserProfileBuilder`.

In this variant, rather than just summing the vectors for all positively-rated items, compute a weighted sum of the item vectors for all rated items, with weights being based on the user's rating. Your solution should implement the following formula:

$$\vec{u} = \sum_{i \in I(u)} (r_{ui} - \mu_u) \vec{i}$$

Example Output for Weighted User Profile

The following example gives actual outputs for users 42 and 91 in the data set. It was executed using `./gradlew recommendWeighted -PuserId=42, 91` in a Unix-like console.

recommendations for user 42:

2081 (Little Mermaid, The (1989)): 0.191
2724 (Runaway Bride (1999)): 0.109
1265 (Groundhog Day (1993)): 0.107
2671 (Notting Hill (1999)): 0.093
96861 (Taken 2 (2012)): 0.082
58299 (Horton Hears a Who! (2008)): 0.077
1202 (Withnail & I (1987)): 0.056
1295 (Unbearable Lightness of Being, The (1988)): 0.047
7361 (Eternal Sunshine of the Spotless Mind (2004)): 0.039
3147 (Green Mile, The (1999)): 0.036

recommendations for user 91:

4816 (Zoolander (2001)): 0.438
8376 (Napoleon Dynamite (2004)): 0.323
1247 (Graduate, The (1967)): 0.296

4361 (Tootsie (1982)): 0.268
6350 (Laputa: Castle in the Sky (Tenkû no shiro Rapyuta) (1986)): 0.216
2167 (Blade (1998)): 0.206
6947 (Master and Commander: The Far Side of the World (2003)): 0.202
5459 (Men in Black II (a.k.a. MIIB) (a.k.a. MIB 2) (2002)): 0.202
2724 (Runaway Bride (1999)): 0.193
2161 (NeverEnding Story, The (1984)): 0.190

Submitting

Submit your code as a zip file on TRACS.

To create this zip file, please use the pre-created archive functionality in the Gradle build:

```
./gradlew prepareSubmission
```

This will ensure that your submission contains all required files. It will produce a submission file in `build/distributions`.